


A Evaluation labels: built from executed evidence, never used as features



B Representations of a failing program: available before any repair



Instructor tool (AAI Lab, Fig. 8): groups with shared evidence, one ILA-2 rule per group, and a mechanism hypothesis only when a rule covers at least half of the group.